

Coding Project N°7 - Music generation: Implement a music generation deep neural network integrating a variational auto-encoder approach with transformers.

Dipartimento di Informatica
Master Degree - Computer Science - AI Curriculum

Intelligent Systems for Pattern Recognition
9 CFU

PROF: Davide Bacciu
Niko Paterniti Barbino 638257



UNIVERSITÀ DI PISA

Index:

1

INTRODUCTION

2

LAKH MIDI
DATASET

3

NEURAL
NETWORK
STRUCTURE

4

EXPERIMENTS

5

RESULTS



INTRODUCTION


The focus of this project is the development of a music generation Deep Neural Network integrating a Variational AutoEncoder approach with Transformers.

For this purpose we used:

- Lakh MIDI dataset to facilitate large-scale music information retrieval
- Pytorch implementation of Compressive Transformers

Lakh MIDI DATASET

The Lakh MIDI dataset is a collection of 176,581 unique MIDI files used to facilitate large-scale music information retrieval, both symbolic (using the MIDI files alone) and audio content-based (using information extracted from the MIDI files as annotations for the matched audio files).



Lakh MIDI Dataset

Data Partition

From the whole “Clean MIDI dataset” we extracted for each song its genre and then we randomly and equally sampled 120 .mid files among 3 genres:

- AlternativeRock
- ArtRock
- BluesRock

Making sure that the sampled songs all had the “acoustic guitar” as instrument (identified by the program ‘25’)

Music21

Music21 is a Python toolkit used for computer-aided musicology. It allows us to teach the fundamentals of music theory, generate music examples and study music. The toolkit provides a simple interface to acquire the musical notation of MIDI files. Additionally, it allows us to create Note and Chord objects so that we can make our own MIDI files easily.

Neural Network Structure

The architecture used in this project is the
“Pytorch Compressive Transformer”:

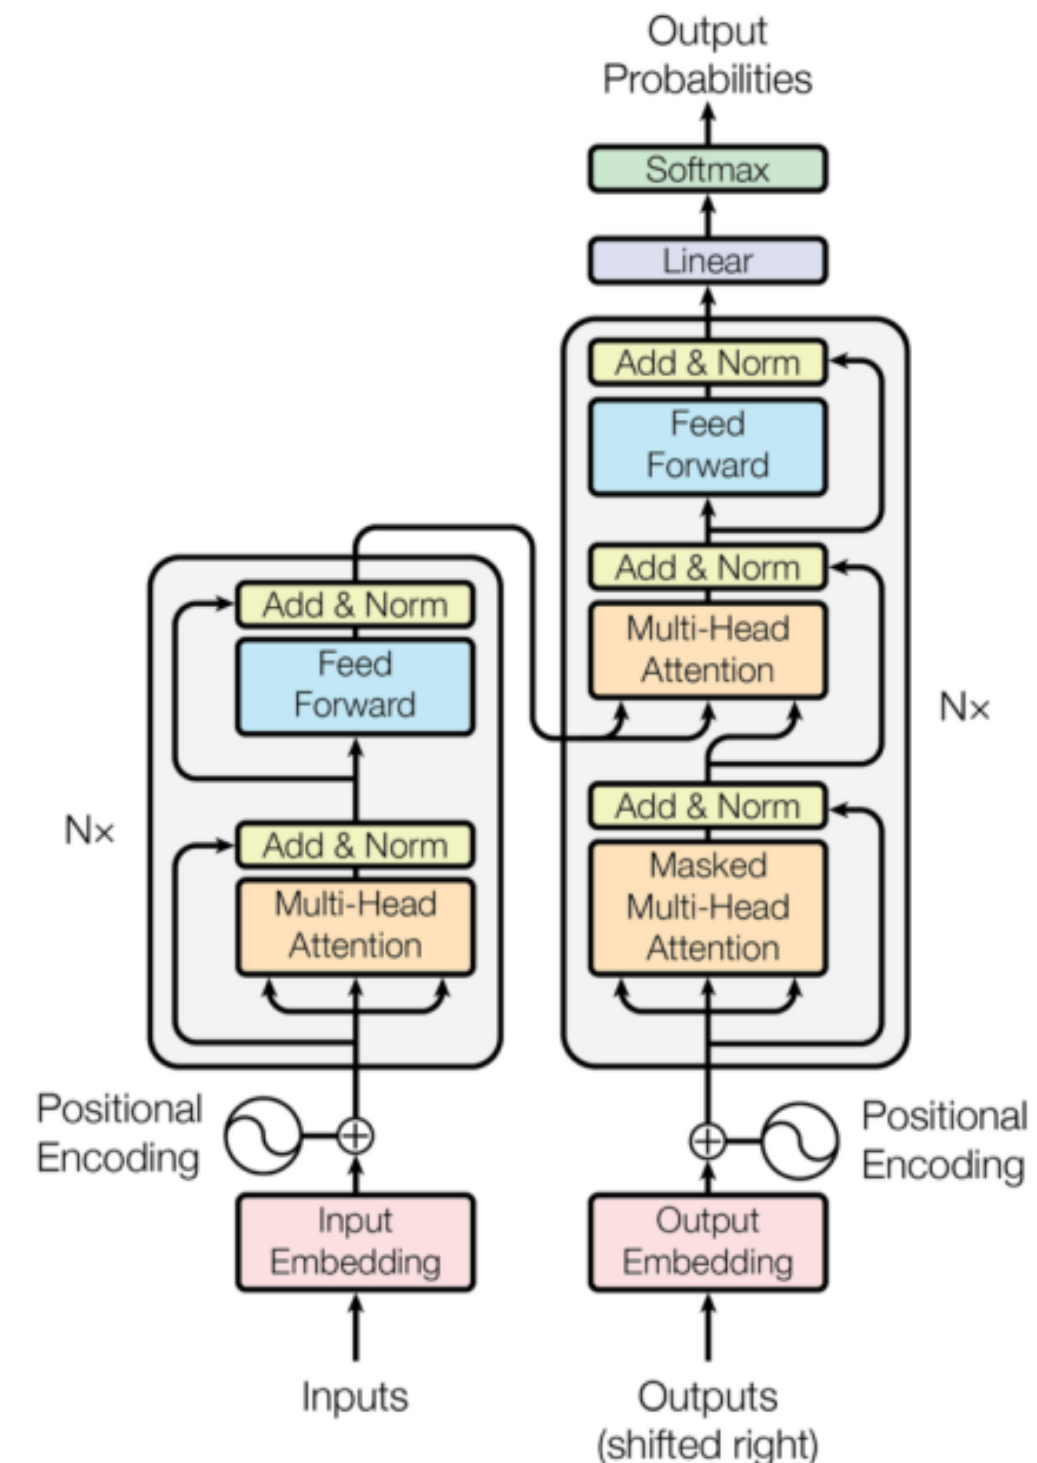
- A variant of Transformer-XL with compressed memory for long-range language modelling.
- An attentive sequence model which compresses past memories for long-range sequence learning.

```
def create_network(sequence_length, n_vocab):  
    model = CompressiveTransformer(  
        num_tokens = n_vocab,  
        dim = sequence_length,  
        depth = 6,  
        seq_len = sequence_length,  
        mem_len = sequence_length,  
        cmem_len = 256,  
        cmem_ratio = 4,  
        memory_layers = [5,6],  
        gru_gated_residual = False,  
        #post-attention dropout  
        attn_dropout = 0.1,  
        #feedforward dropout  
        ff_dropout = 0.1,  
        #attention layer output dropout  
        attn_layer_dropout = 0.1  
    )  
  
    model = AutoregressiveWrapper(model)  
    model.cuda()  
    return model
```


Neural Network Structure

Compressive Transformer

- Based on the Transformer XL architecture with the addition of compressed memory.
- Encoder-decoder design
- Self-Attention
- Data flow: given a sequence of tokens x they are mapped in a continuous representation z . The decoder takes as input one element of z at a time with the additional input of the previous output generated by itself.
- AutoRegressive Model



Neural Network Structure

Self Attention:

- Learn and describe the relationship between the sequence tokens.
- Have an attention function that given a vector query Q , the matrices of keys K and values V (dimension d^k) computes their relationship.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{Q * K^T}{\sqrt{d^k}}\right) * V$$

Neural Network Structure

Recurrence Mechanism

Technique that enables long-term dependencies using information from previous segments conditioning the values of K and the V of the segment with the previously hidden state sequence produced by the previous segment concatenated with the earlier hidden state sequence produced by the current segment.


Relative Positional Encoding

Introduces a relative position encoding to correctly compute the self attention of the current segment because the current sequence has the same positional encoding as the previous

Neural Network Structure

Memory Compression

The Recurrence Mechanisms leads us to the requirement of additional memory to store the informations of past segments. The memory compression helps us store as much information as we can through the compressed memory.



Example of compressed memory in a transformer:

- 3 layer,
- sequence length = 3,
- memory size = 6,
- compressed memory size = 6
- rate of compression = 3.

EXPERIMENTS

To run the code and perform the experiments i bought **Google Colab Pro** due to the long time required for the training and the limited resources available with the free version.

So the training was computed on Google Colab using the **Tesla V100 machine** (See the image on the left for more details)

```

+-----+
| NVIDIA-SMI 525.105.17    Driver Version: 525.105.17    CUDA Version: 12.0
```

GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute	M. MIG M.
0	Tesla V100-SXM2...	Off	00000000:00:04.0	Off			0
N/A	39C	P0	25W / 300W	0MiB / 16384MiB	0%	Default	N/A

```

+-----+

+-----+
| Processes:
```

GPU	GI	CI	PID	Type	Process name	GPU Memory
	ID	ID				Usage
No running processes found						

```

+-----+

```

Experiments

- **Architecture:** based on the implementation of the following notebook:
<https://github.com/lucidrains/compressive-transformer-pytorch>
- **Input:** Vector of 64 notes and chords.
- **6 layers** for the encoder and the decoder
- **Optimizer:** Adam PyTorch
- **Sequence length:** 32, 64, 128. Due to time required for training, probability of overfitting/underfitting and other criteria i decided to use 64 as sequence length.
- **Regularization:** Dropout = 0.1

I studied the task and the strategy used in this project from the paper:

<https://towardsdatascience.com/how-to-generate-music-using-a-lstm-neural-network-in-keras-68786834d4c5>

<https://colinraffel.com/publications/thesis.pdf>

And the relative code demonstration to better understand the task of music generation:

<https://colab.research.google.com/drive/19TQqekOlnOSW36VCL8CPVEQKBBukmaEQ#scrollTo=DDOBVWULXfpz>

Experiments

Initially i performed my experiments with the whole dataset of 120 MIDI files equally and randomly divided among 3 genres (AlternativeRock, ArtRock, BluesRock) setting a very high number of epochs. I found out that the time required for each epoch was in the range of 500 seconds so the time to get the results was definitely too high and i quickly run out of resources from the Colab runtime, so i decided to perform the training on a smaller portion of the dataset using a total of 40 MIDI files for the training and the validation with a time for each epoch in the range of 25 seconds

Results

Epoch: 0 Training loss: 0.8095 Epoch: 1 Training loss: 0.8072 Epoch: 2 Training loss: 0.8051 Epoch: 3 Training loss: 0.8027 Epoch: 4 Training loss: 0.8011 validation loss: 6.0538	Epoch: 20 Training loss: 0.7668 Epoch: 21 Training loss: 0.7654 Epoch: 22 Training loss: 0.7622 Epoch: 23 Training loss: 0.7618 Epoch: 24 Training loss: 0.7585 validation loss: 6.1651	Epoch: 40 Training loss: 0.7258 Epoch: 41 Training loss: 0.7243 Epoch: 42 Training loss: 0.7219 Epoch: 43 Training loss: 0.7199 Epoch: 44 Training loss: 0.7173 validation loss: 6.2856	Epoch: 40 Training loss: 0.7258 Epoch: 41 Training loss: 0.7243 Epoch: 42 Training loss: 0.7219 Epoch: 43 Training loss: 0.7199 Epoch: 44 Training loss: 0.7173 validation loss: 6.2856	Epoch: 60 Training loss: 0.6856 Epoch: 61 Training loss: 0.6828 Epoch: 62 Training loss: 0.6808 Epoch: 63 Training loss: 0.6787 Epoch: 64 Training loss: 0.6771 validation loss: 6.4080
Epoch: 5 Training loss: 0.7980 Epoch: 6 Training loss: 0.7956 Epoch: 7 Training loss: 0.7951 Epoch: 8 Training loss: 0.7916 Epoch: 9 Training loss: 0.7905 validation loss: 6.0906	Epoch: 25 Training loss: 0.7577 Epoch: 26 Training loss: 0.7546 Epoch: 27 Training loss: 0.7529 Epoch: 28 Training loss: 0.7507 Epoch: 29 Training loss: 0.7482 validation loss: 6.2097	Epoch: 45 Training loss: 0.7164 Epoch: 46 Training loss: 0.7135 Epoch: 47 Training loss: 0.7116 Epoch: 48 Training loss: 0.7091 Epoch: 49 Training loss: 0.7081 validation loss: 6.3383	Epoch: 45 Training loss: 0.7164 Epoch: 46 Training loss: 0.7135 Epoch: 47 Training loss: 0.7116 Epoch: 48 Training loss: 0.7091 Epoch: 49 Training loss: 0.7081 validation loss: 6.3383	Epoch: 65 Training loss: 0.6751 Epoch: 66 Training loss: 0.6731 Epoch: 67 Training loss: 0.6716 Epoch: 68 Training loss: 0.6684 Epoch: 69 Training loss: 0.6671 validation loss: 6.4736
Epoch: 10 Training loss: 0.7874 Epoch: 11 Training loss: 0.7859 Epoch: 12 Training loss: 0.7829 Epoch: 13 Training loss: 0.7819 Epoch: 14 Training loss: 0.7792 validation loss: 6.1258	Epoch: 30 Training loss: 0.7462 Epoch: 31 Training loss: 0.7451 Epoch: 32 Training loss: 0.7418 Epoch: 33 Training loss: 0.7404 Epoch: 34 Training loss: 0.7374 validation loss: 6.2456	Epoch: 50 Training loss: 0.7050 Epoch: 51 Training loss: 0.7041 Epoch: 52 Training loss: 0.7011 Epoch: 53 Training loss: 0.6995 Epoch: 54 Training loss: 0.6967 validation loss: 6.3687	Epoch: 50 Training loss: 0.7050 Epoch: 51 Training loss: 0.7041 Epoch: 52 Training loss: 0.7011 Epoch: 53 Training loss: 0.6995 Epoch: 54 Training loss: 0.6967 validation loss: 6.3687	Epoch: 70 Training loss: 0.6648 Epoch: 71 Training loss: 0.6634 Epoch: 72 Training loss: 0.6608 Epoch: 73 Training loss: 0.6593 Epoch: 74 Training loss: 0.6565 validation loss: 6.4970
Epoch: 15 Training loss: 0.7785 Epoch: 16 Training loss: 0.7751 Epoch: 17 Training loss: 0.7733 Epoch: 18 Training loss: 0.7707 Epoch: 19 Training loss: 0.7693 validation loss: 6.1418	Epoch: 35 Training loss: 0.7360 Epoch: 36 Training loss: 0.7338 Epoch: 37 Training loss: 0.7326 Epoch: 38 Training loss: 0.7298 Epoch: 39 Training loss: 0.7286 validation loss: 6.2593	Epoch: 55 Training loss: 0.6952 Epoch: 56 Training loss: 0.6933 Epoch: 57 Training loss: 0.6908 Epoch: 58 Training loss: 0.6886 Epoch: 59 Training loss: 0.6873 validation loss: 6.3766	Epoch: 55 Training loss: 0.6952 Epoch: 56 Training loss: 0.6933 Epoch: 57 Training loss: 0.6908 Epoch: 58 Training loss: 0.6886 Epoch: 59 Training loss: 0.6873 validation loss: 6.3766	Epoch: 75 Training loss: 0.6547 Epoch: 76 Training loss: 0.6532 Epoch: 77 Training loss: 0.6517 Epoch: 78 Training loss: 0.6501 Epoch: 79 Training loss: 0.6468 validation loss: 6.4955

CONCLUSION

In this study we implemented a deep neural network for music generation integrating a variational auto-encoder approach with Transformers, I took inspiration from different projects and papers where I found many informations and different ways of approaching the problem such as LSTM, BI-LSTM ecc... Due to the self attention mechanism the best way to approach this problem is transformer based approach, still it would be interesting making comparisons and trying to obtain better quality songs with a bigger dataset. Also we could study the behaviour of the model considering the encoding of more than one instrument for each note with also the temporal parameter set, so that more than one instrument could be associated with each note or chord generated making music more harmonious

Thanks



UNIVERSITÀ DI PISA

Niko Paterniti Barbino
638257